

```
In [1]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler, OneHotEncoder, OrdinalEncoder
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.linear_model import LinearRegression, Ridge, Lasso, RidgeCV
from sklearn.metrics import mean_squared_error, mean_squared_log_error
import seaborn as sns # Import seaborn for better visualizations
import math
from scipy import stats
from sklearn.ensemble import RandomForestRegressor, BaggingRegressor, GradientBoostingRegressor
from sklearn.svm import SVR
from sklearn.neighbors import KNeighborsRegressor
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
In [2]: from utils import line_plot_viz, box_plot_viz, heat_map_viz, hist_plot_viz,
import warnings
warnings.filterwarnings('ignore')
```

Context This classic dataset contains the prices and other attributes of almost 54,000 diamonds. It's a great dataset for beginners learning to work with data analysis and visualization.

Content price price in US dollars (326 — 18,823)

carat weight of the diamond (0.2--5.01)

cut quality of the cut (Fair, Good, Very Good, Premium, Ideal)

color diamond colour, from J (worst) to D (best)

clarity a measurement of how clear the diamond is (I1 (worst), SI2, SI1, VS2, VS1, VVS2, VVS1, IF (best))

x length in mm (0--10.74)

y width in mm (0--58.9)

z depth in mm (0--31.8)

depth total depth percentage = $z / \text{mean}(x, y) = 2 * z / (x + y)$ (43--79)

table width of top of diamond relative to widest point (43--95)

```
In [3]: og_data = pd.read_csv('./data/Diamonds/diamonds.csv')
data = og_data.copy()
```

```
In [4]: data.describe()
```

Out [4]:

	index	carat	depth	table	price
count	53940.000000	53940.000000	53940.000000	53940.000000	53940.000000
mean	26970.500000	0.797940	61.749405	57.457184	3932.799722
std	15571.281097	0.474011	1.432621	2.234491	3989.439738
min	1.000000	0.200000	43.000000	43.000000	326.000000
25%	13485.750000	0.400000	61.000000	56.000000	950.000000
50%	26970.500000	0.700000	61.800000	57.000000	2401.000000
75%	40455.250000	1.040000	62.500000	59.000000	5324.250000
max	53940.000000	5.010000	79.000000	95.000000	18823.000000

In [5]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 53940 entries, 0 to 53939
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   index      53940 non-null  int64
1   carat      53940 non-null  float64
2   cut        53940 non-null  object
3   color      53940 non-null  object
4   clarity    53940 non-null  object
5   depth      53940 non-null  float64
6   table      53940 non-null  float64
7   price      53940 non-null  int64
8   x          53940 non-null  float64
9   y          53940 non-null  float64
10  z          53940 non-null  float64
dtypes: float64(6), int64(2), object(3)
memory usage: 4.5+ MB
```

In [6]: `data.columns`

```
Out[6]: Index(['index', 'carat', 'cut', 'color', 'clarity', 'depth', 'table', 'price',
              'x', 'y', 'z'],
              dtype='object')
```

Check Null/missing Values

In [7]: `data.isnull().sum()`

```
Out[7]: index      0
        carat      0
        cut        0
        color      0
        clarity    0
        depth      0
        table      0
        price      0
        x          0
        y          0
        z          0
        dtype: int64
```

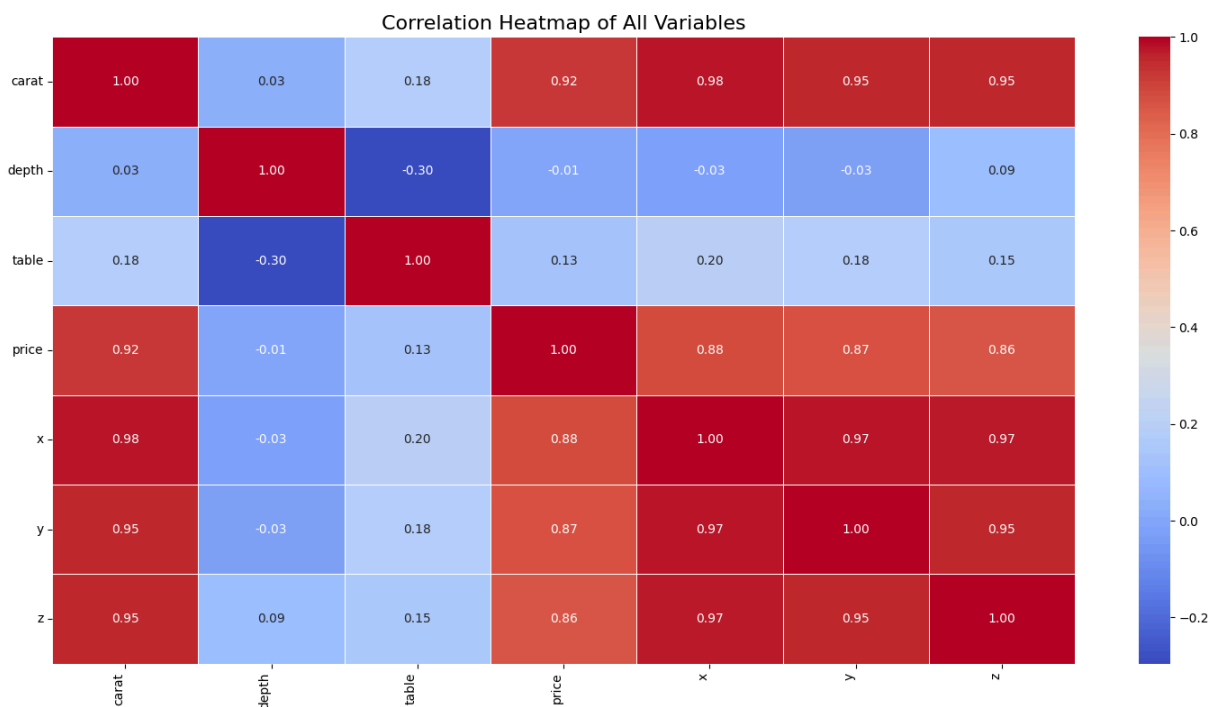
Remove columns

```
In [8]: # Drop index column
data = data.drop('index', axis=1)
```

```
In [9]: data.dtypes[0]
```

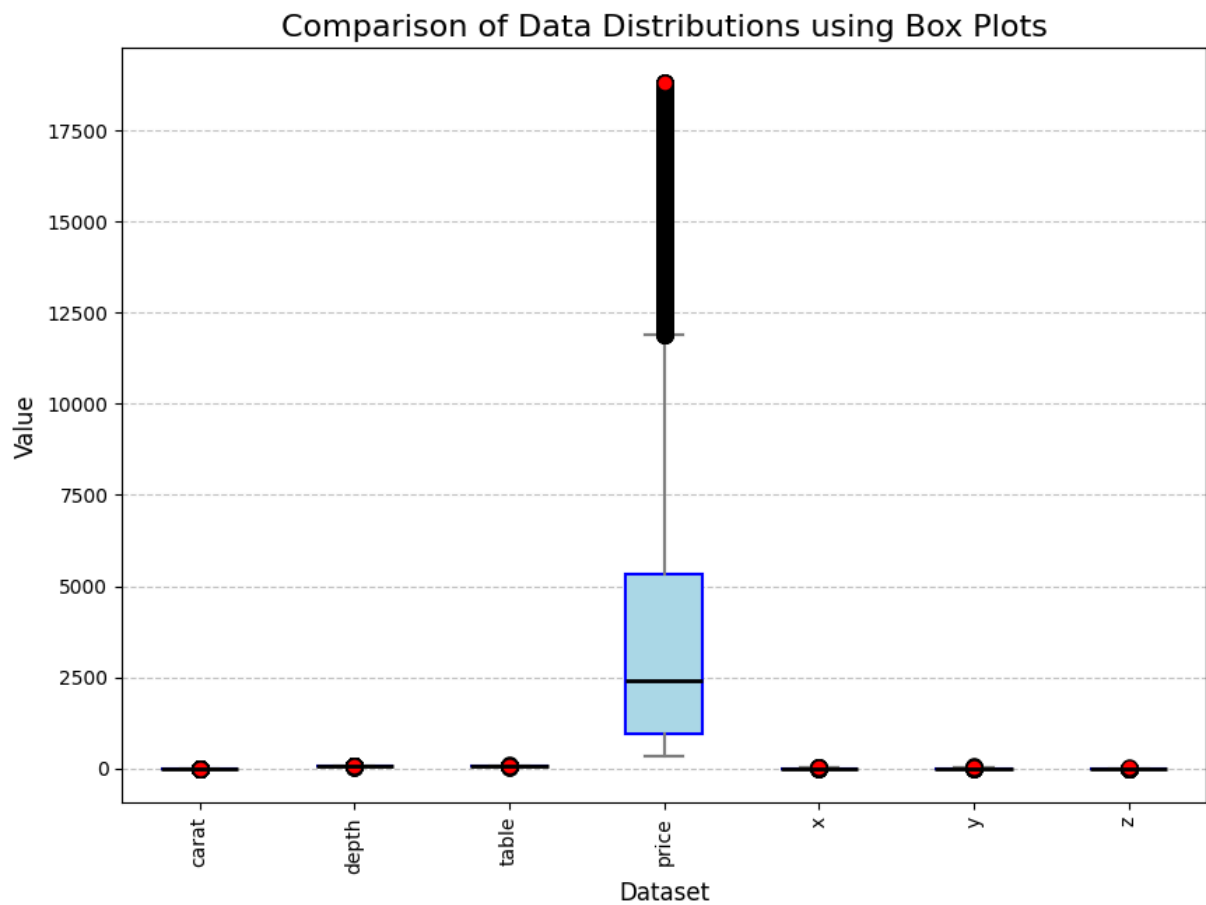
```
Out[9]: dtype('float64')
```

```
In [10]: cols= list(data.columns)
numeric_data = data.select_dtypes(include=['float64', 'int64'])
numeric_cols = list(numeric_data.columns)
heat_map_viz(numeric_data)
```

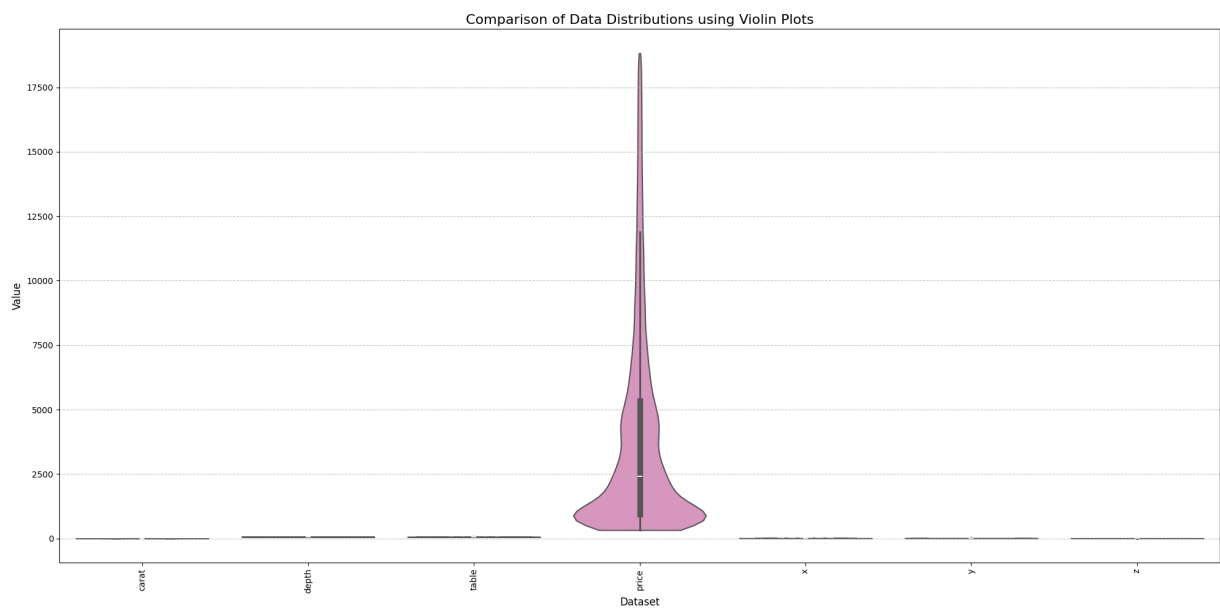


```
In [11]: box_plot_viz(numeric_data)
```

<Figure size 2000x1000 with 0 Axes>




```
In [12]: violin_plot_viz(numeric_data)
```



Scale data for the column price


```
In [13]: scaler = StandardScaler()
numeric_data['price'] = scaler.fit_transform(numeric_data[['price']])
numeric_data['depth'] = scaler.fit_transform(numeric_data[['depth']])
```

```
numeric_data['table'] = scaler.fit_transform(numeric_data[['table']])
violin_plot_viz(numeric_data)
```



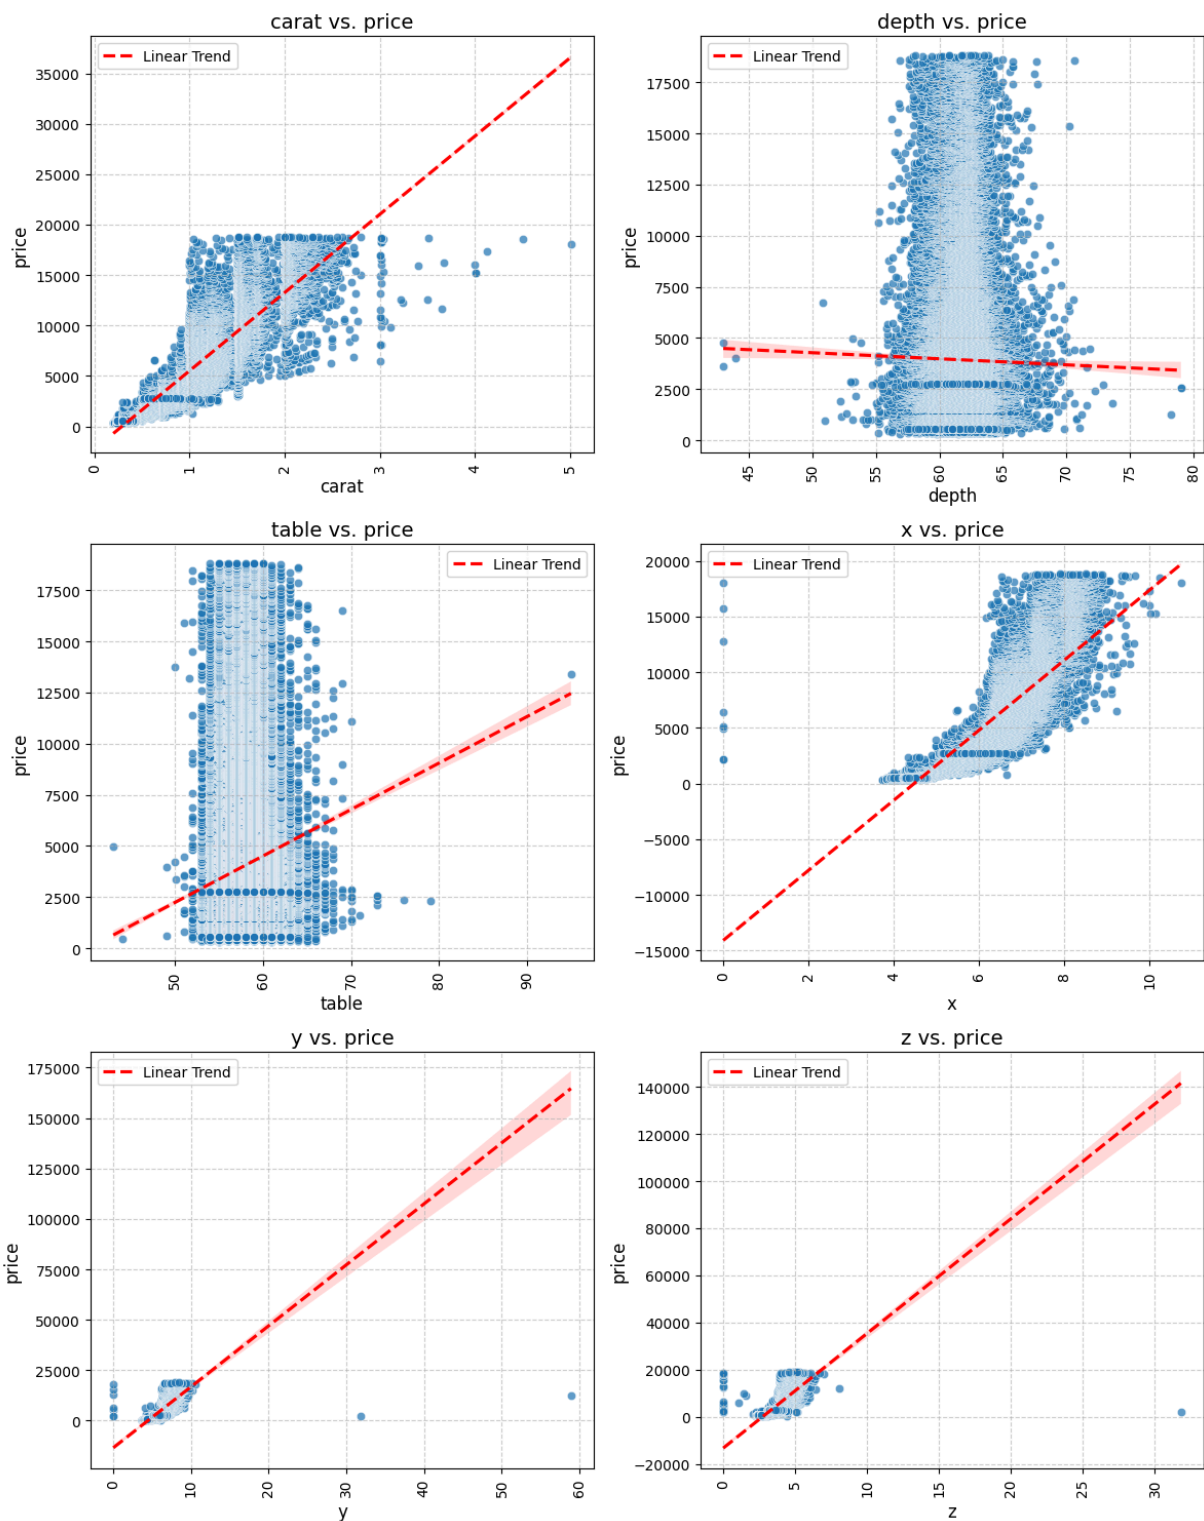
```
In [14]: box_plot_viz(numeric_data)
```

<Figure size 2000x1000 with 0 Axes>

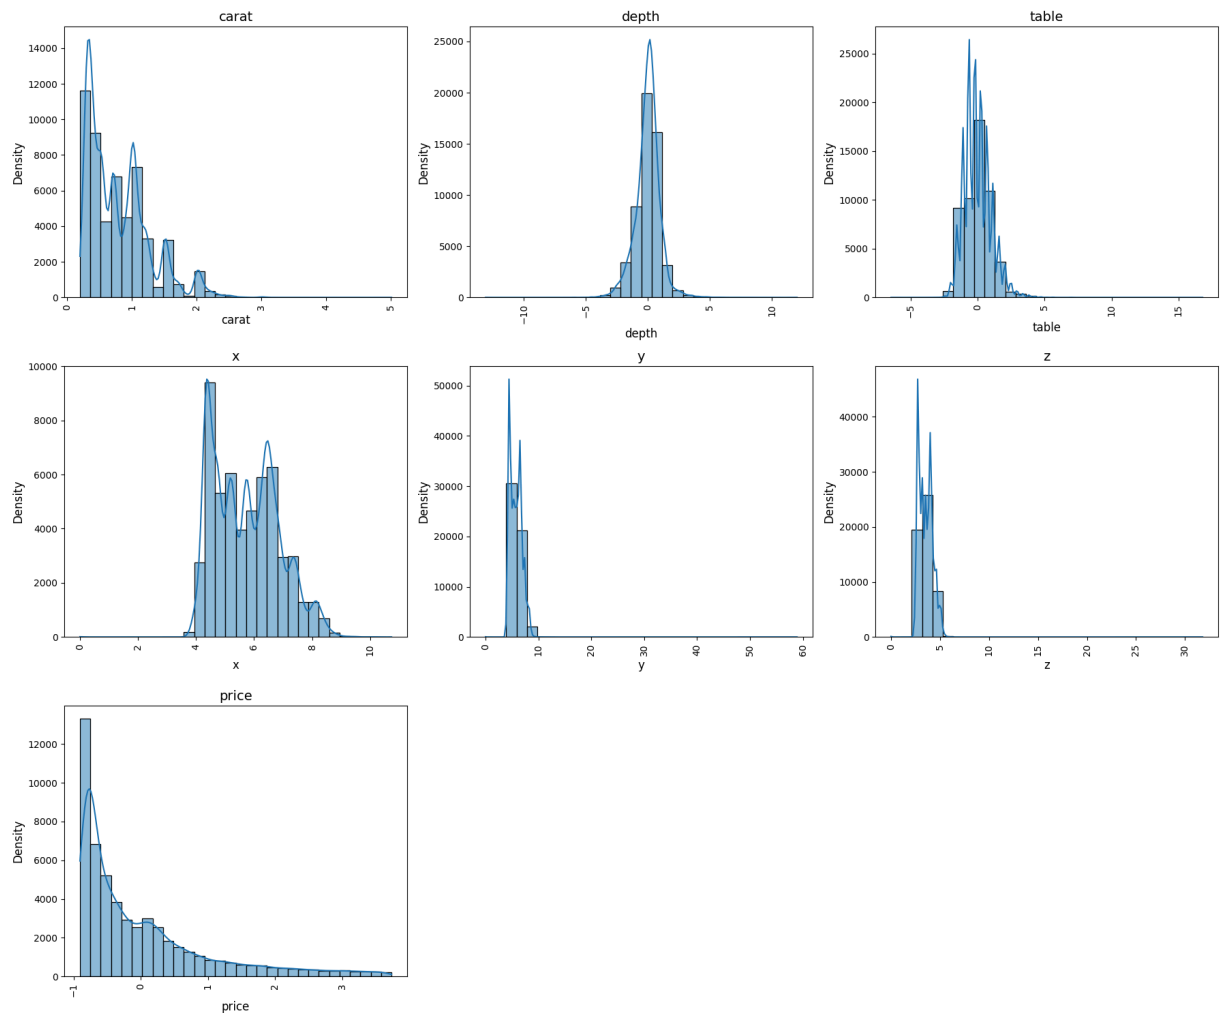


```
In [15]: X_cols = ['carat', 'depth', 'table', 'x', 'y', 'z']
Y_cols = ['price']
```

```
line_plot_viz(X_columns=X_cols, Y_columns=Y_cols, data=data)
```



```
In [16]: hist_plot_viz(columns=[*X_cols, *Y_cols], data=numeric_data)
```



VIZ Analysis to test for multi-collinearity

- From the heat map we can see that X, Y, Z and carat have high multicollinearity

```
In [17]: # VIF calculation
X_scaled = StandardScaler().fit_transform(data[numeric_cols].drop(['price'],
vif_data = pd.DataFrame()
vif_data['feature'] = data[numeric_cols].drop('price', axis=1).columns
vif_data['VIF'] = [variance_inflation_factor(X_scaled, i) for i in range(X_s
vif_data.sort_values('VIF', ascending=False)
```

```
Out[17]:
```

	feature	VIF
3	x	56.187704
5	z	23.530049
0	carat	21.602712
4	y	20.454295
1	depth	1.496590
2	table	1.143225

```
In [18]: # VIF calculation after removing highly multi-collinear features
X_scaled = StandardScaler().fit_transform(data[numeric_cols].drop(['price',
vif_data = pd.DataFrame()
vif_data['feature'] = data[numeric_cols].drop(['price', 'x', 'y', 'z'], axis
vif_data['VIF'] = [variance_inflation_factor(X_scaled, i) for i in range(X_s
vif_data.sort_values('VIF', ascending=False)
```

```
Out[18]:
```

	feature	VIF
2	table	1.141032
1	depth	1.104275
0	carat	1.042039

Feature Removal

- Remove x,y,z as they have high multicollinearity with carat

```
In [19]: cols_to_remove = ['x', 'y', 'z']
numeric_data.drop(columns= cols_to_remove, inplace=True)
proccssd_numeric_cols = [x for x in X_cols if x not in cols_to_remove]
proccssd_numeric_cols
```

```
Out[19]: ['carat', 'depth', 'table']
```

Identify the categorical variables

```
In [20]: def convert_to_categorical_dtypes(cat_threshold, cols, proc_data):
data = proc_data.copy()
for feature in cols:
    uniq_vals = data[feature].unique()
    if (len(uniq_vals) < cat_threshold):
        data[feature] = data[feature].astype('category')
print(data[cols].dtypes)
return data
cat_threshold = 30
data = convert_to_categorical_dtypes(cat_threshold=cat_threshold, cols=data.
```

```
carat      float64
cut        category
color      category
clarity     category
depth      float64
table      float64
price      int64
x          float64
y          float64
z          float64
dtype: object
```

For the Categorical variables lets aggregate demand by each category and plot


```
In [21]: cat_cols = data.select_dtypes(include='category').columns
num_features = len(cat_cols)
n_cols = 3 # You can adjust the number of columns in the grid
n_rows = (num_features + n_cols - 1) // n_cols # Calculate rows needed
colors = ['r', 'g', 'b', 'm', 'c']
fig, axes = plt.subplots(n_rows, n_cols, figsize=(n_cols * 6, n_rows * 5))
axes = axes.flatten() # Flatten the 2D array of axes for easy iteration
i = 0

for feature in cat_cols:
    uniq_vals = data[feature].unique()
    cat_average = data.groupby(feature)[*Y_cols].mean()


    # 3. Iterate through each feature and plot its relationship with the target
    ax = axes[i] # Get the current subplot axis

    sns.barplot(data=cat_average, x= feature, y= 'price', hue=feature, ax=ax,

    ax.set_title(f'{feature}', fontsize=14)
    ax.set_xlabel(feature, fontsize=12)
    ax.set_ylabel('Demand', fontsize=12)
    ax.set_xticklabels(ax.get_xticklabels(), rotation=90)

    i = i + 1
for j in range(i, len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout() # Adjust layout to prevent overlapping titles/labels
plt.show()
```



Remove columns

- No categorical columns to remove

Create the processed data by combining the processed numeric and categorical columns

```
In [22]: processed_data = pd.concat([data[cat_cols], numeric_data], axis=1)
processed_data.head()
```

Out [22]:

	cut	color	clarity	carat	depth	table	price
0	Ideal	E	SI2	0.23	-0.174092	-1.099672	-0.904095
1	Premium	E	SI1	0.21	-1.360738	1.585529	-0.904095
2	Good	E	VS1	0.23	-3.385019	3.375663	-0.903844
3	Premium	I	VS2	0.29	0.454133	0.242928	-0.902090
4	Good	J	SI2	0.31	1.082358	0.242928	-0.901839

In [23]: `processed_data.dtypes`

Out [23]:

```
cut          category
color        category
clarity      category
carat        float64
depth        float64
table        float64
price        float64
dtype: object
```

Encode the categorical variables using the below

cut quality of the cut (Fair, Good, Very Good, Premium, Ideal)

color diamond colour, from J (worst) to D (best)

clarity a measurement of how clear the diamond is (I1 (worst), SI2, SI1, VS2, VS1, VVS2, VVS1, IF (best))

```
In [24]: encoder = OrdinalEncoder(categories=[['Fair', 'Good', 'Very Good', 'Premium', 'Ideal']])
processed_data['cut'] = encoder.fit_transform(processed_data[['cut']])

encoder = OrdinalEncoder(categories=[['J', 'I', 'H', 'G', 'F', 'E', 'D']])
processed_data['color'] = encoder.fit_transform(processed_data[['color']])


encoder = OrdinalEncoder(categories=[['I1', 'SI2', 'SI1', 'VS2', 'VS1', 'VVS2', 'VVS1', 'IF']])
processed_data['clarity'] = encoder.fit_transform(processed_data[['clarity']])

processed_data.head()
```

Out [24]:

	cut	color	clarity	carat	depth	table	price
0	4.0	5.0	1.0	0.23	-0.174092	-1.099672	-0.904095
1	3.0	5.0	2.0	0.21	-1.360738	1.585529	-0.904095
2	1.0	5.0	4.0	0.23	-3.385019	3.375663	-0.903844
3	3.0	1.0	3.0	0.29	0.454133	0.242928	-0.902090
4	1.0	0.0	1.0	0.31	1.082358	0.242928	-0.901839

```
In [25]: heat_map_viz(processed_data)
```



```
In [26]: Y_cols = ['price']
Y = processed_data[Y_cols]
X = processed_data.drop(Y_cols, axis=1)
X_cols = list(X.columns)
```

```
print('Target column is: ', Y_cols)
print('Features are: ', X_cols)
```


Target column is: ['price']

Features are: ['cut', 'color', 'clarity', 'carat', 'depth', 'table']

A look at the distribution of the target variable Price


- According to the plot below the distribution is not normal so let us apply a log normal transform

```
In [27]: hist_plot_viz(Y_cols, processed_data)
```




Below is the graph after the log normal transformation

```
In [28]: processed_data[Y_cols] = np.log(data[Y_cols])  
hist_plot_viz(Y_cols, processed_data)
```



```
In [29]: sns.histplot(data=processed_data, x='price', hue='cut', kde=True, multiple='')
```

```
Out[29]: <Axes: xlabel='price', ylabel='Count'>
```



```
In [30]: df = processed_data.copy()
df['price_per_carat'] = df['price'] / df['carat']
df.groupby('cut')['price_per_carat'].mean().sort_values()
```

```
Out[30]: cut
0.0      9.226632
1.0     11.575319
3.0     11.755446
2.0     12.499832
4.0     13.902481
Name: price_per_carat, dtype: float64
```

```
In [31]: df.groupby('cut')['price'].mean().sort_values()
```

```
Out[31]: cut
4.0      7.639467
2.0      7.798664
1.0      7.842809
3.0      7.950795
0.0      8.093441
Name: price, dtype: float64
```

```
In [32]: df = processed_data.copy()
df['price_per_carat'] = df['price'] / df['carat']
df.groupby('color')['price_per_carat'].mean().sort_values()
```

```
Out[32]: color
0.0      8.984839
1.0     10.415987
2.0     11.514335
3.0     12.933443
4.0     13.152478
6.0     14.149139
5.0     14.256642
Name: price_per_carat, dtype: float64
```

```
In [33]: df = processed_data.copy()
df['price_per_carat'] = df['price'] / df['carat']
df.groupby('clarity')['price_per_carat'].mean().sort_values()
```

```
Out[33]: clarity
0.0      7.419620
1.0      9.182076
2.0     11.441858
3.0     13.001680
4.0     13.477317
5.0     15.831236
6.0     17.426974
7.0     17.652448
Name: price_per_carat, dtype: float64
```

```
In [34]: df['cut'].value_counts().sort_index()
```

```
Out[34]: cut
0.0      1610
1.0      4906
2.0     12082
3.0     13791
4.0     21551
Name: count, dtype: int64
```


Some important observations from above results

- Though the data legend suggests that there is ordinality in the cut feature ie better cut should be priced higher but if you look at the effect of each type of cut on the price. It does not increase as the cuts get better. Infact I have a hypothesis that the mid quality cuts are overpriced and earn the most margin for the seller. There is enough data for the cut type 2 when compared to the better cuts so imbalance in the classes can also be ruled out.
- We need to remove the ordinality encoding and add one hot encoding.
- All the other categorical variables do show a trend in price due to ordinality so we will preserve that.

```
In [35]: cuts_df = pd.get_dummies(data['cut'], drop_first=True, dtype='int')
processed_data = pd.concat([processed_data, cuts_df], axis=1)
```

```
In [36]: processed_data.drop('cut', inplace=True, axis=1)
```

In [37]: `heat_map_viz(processed_data)`



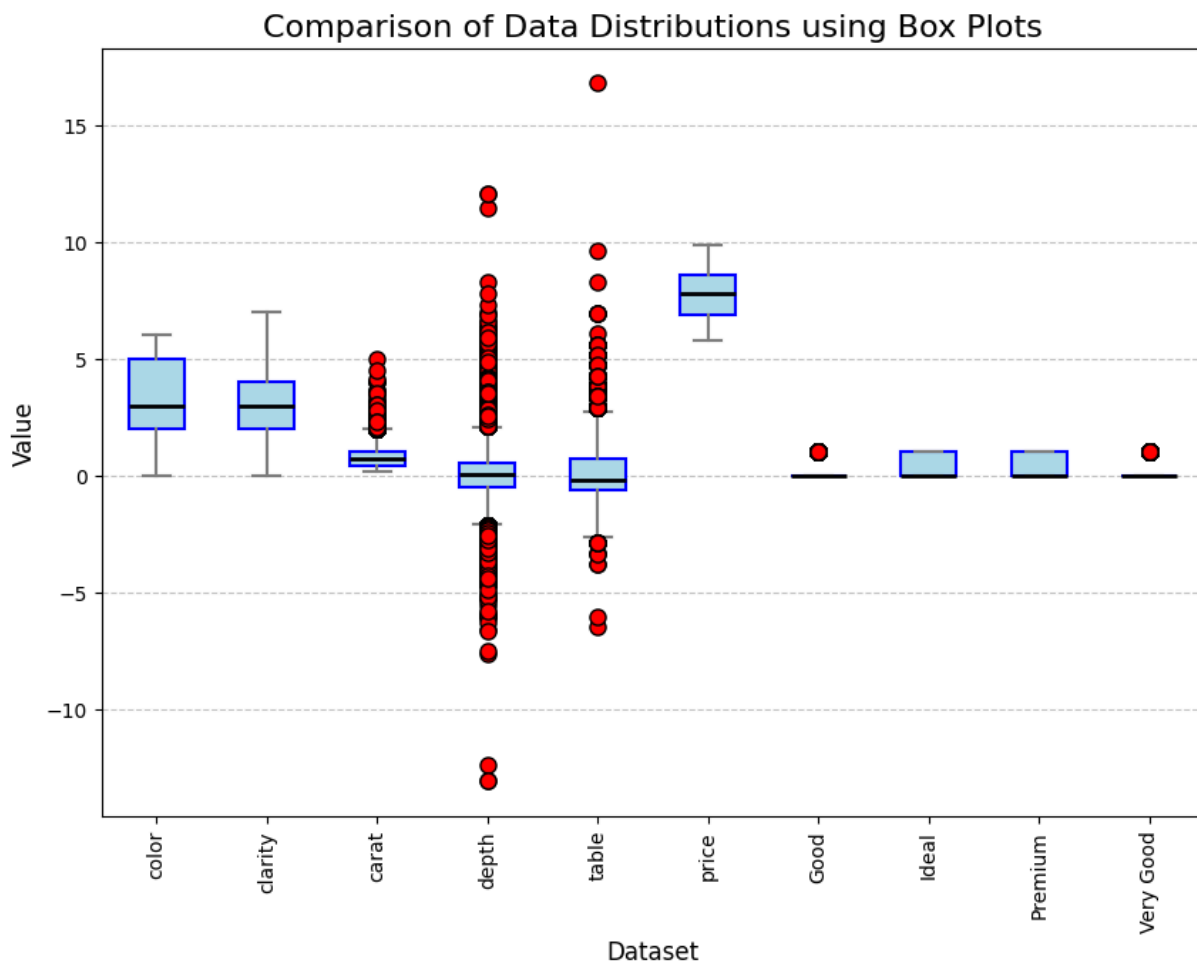
In [38]: `processed_data.dtypes`

```
Out[38]: color      float64
clarity    float64
carat      float64
depth      float64
table      float64
price      float64
Good       int64
Ideal      int64
Premium    int64
Very Good  int64
dtype: object
```

Detect outliers If any

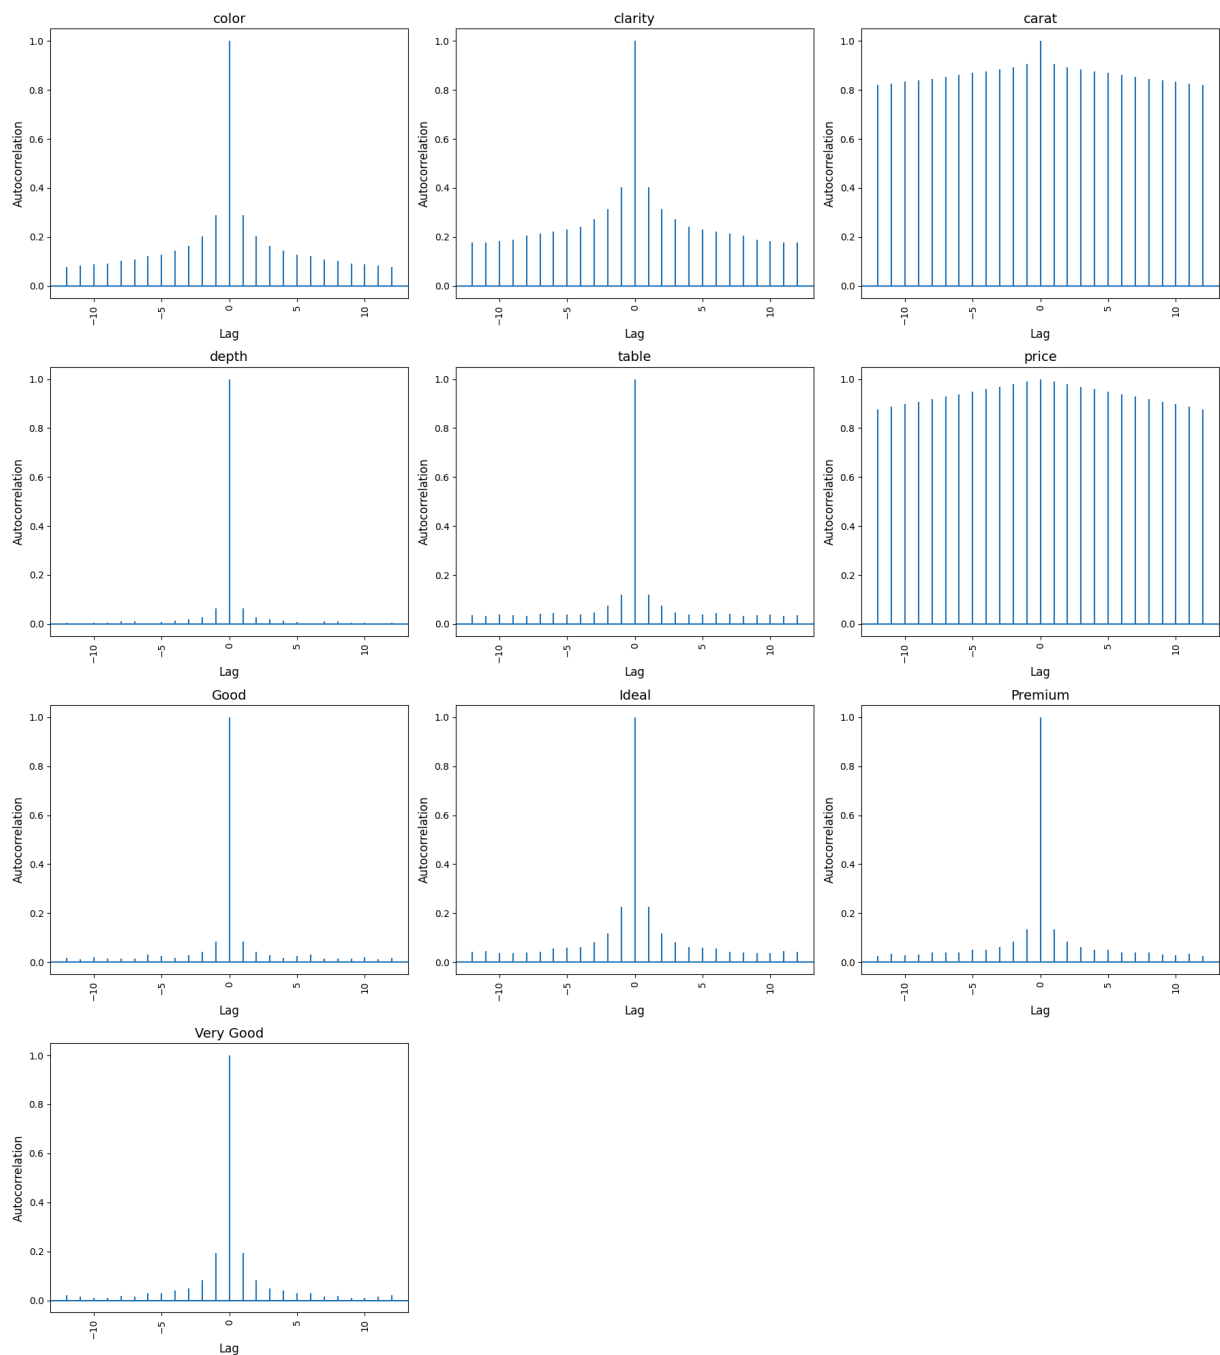
In [39]: `box_plot_viz(processed_data)`

<Figure size 2000x1000 with 0 Axes>



Test for autocorrelation in the target column

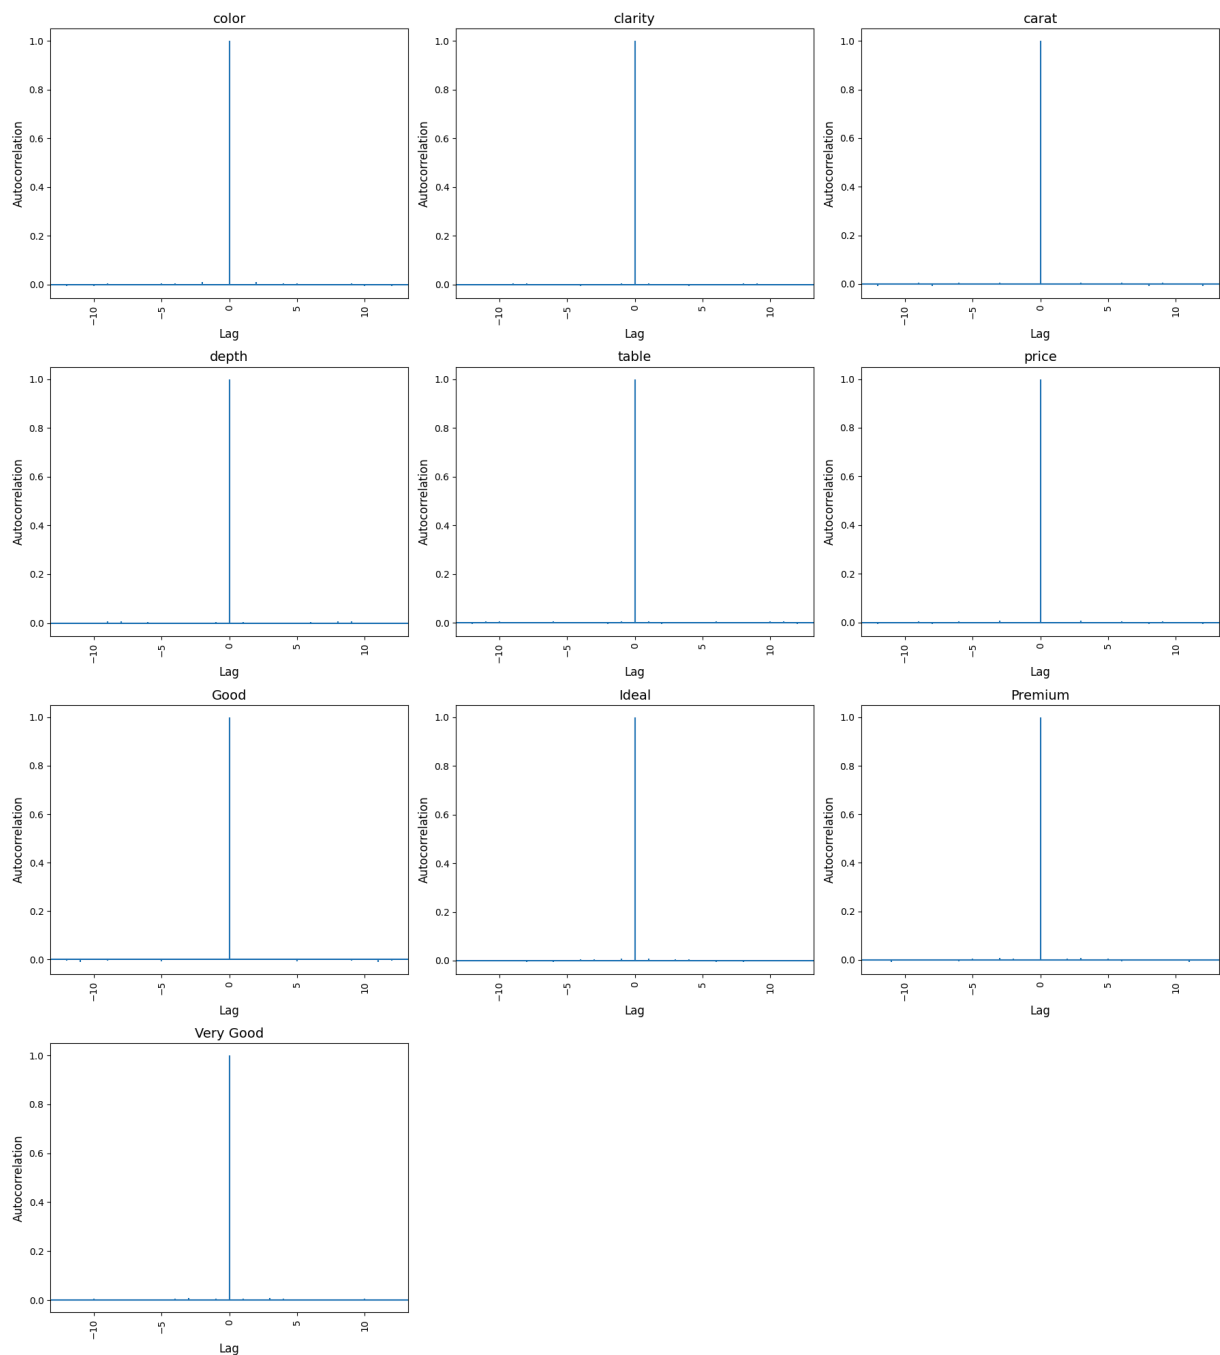
```
In [40]: auto_corr_viz(processed_data=processed_data)
```



Looks like there is high auto correlation in the features

- lets reshuffle the data and test this again

```
In [41]: df_shuffled = processed_data.sample(frac=1, random_state=42).reset_index(drop=True)
         auto_corr_viz(df_shuffled)
```



No auto correlation found

Train Test Split

```
In [42]: X = processed_data.drop(columns=['price'], axis=1)
y = processed_data['price']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

Fit the linear regression model to get the baseline

```
In [43]: model = LinearRegression()
model.fit(X = X_train, y = y_train)
```

```
Out[43]:
```

▼ LinearRegression ⓘ ?

► Parameters

R Squared metrics of the model

```
In [44]: r2_train = model.score(X = X_train, y = y_train)
print('r2 train : ', r2_train)
r2_test = model.score(X = X_test, y = y_test)
print('r2 test : ', r2_test)
```

```
r2 train : 0.8807282361735038
r2 test : 0.8786308335563693
```

```
In [45]: # RMSE
y_pred = model.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print(f"RMSE: {rmse:.4f}")
```

```
RMSE: 0.3538
```

```
In [46]: rmsle_lin = np.sqrt(mean_squared_log_error(y_test, y_pred))
print(f"RMSLE: {rmsle_lin:.4f}")
```

```
RMSLE: 0.0387
```

Try other models

```
In [47]: models=[RandomForestRegressor(),AdaBoostRegressor(),BaggingRegressor(),SVR()]
model_names=['RandomForestRegressor','AdaBoostRegressor','BaggingRegressor',
rmsle=[]
d={}
for model in range (len(models)):
    clf=models[model]
    clf.fit(X_train,y_train)
    test_pred=clf.predict(X_test)
    rmsle.append(np.sqrt(mean_squared_log_error(y_test, test_pred)))
d={'Modelling Algo':model_names,'RMSLE':rmsle}
```

```
In [48]: model_metrics = pd.DataFrame(d)
model_metrics.loc[len(model_metrics)] = {'Modelling Algo':'Linear regressior'
model_metrics.sort_values(by='RMSLE')
```

Out [48]:

	Modelling Algo	RMSLE
0	RandomForestRegressor	0.013249
2	BaggingRegressor	0.013616
3	SVR	0.015268
1	AdaBoostRegressor	0.023683
5	Linear regression	0.038740
4	KNeighborsRegressor	0.040098

Looks like Random forest is our best bet

- Lets do some cross validation to confirm this.

```
In [49]: scores = cross_val_score(RandomForestRegressor(), X, y, cv=5, scoring='neg_m
rmsle_scores = np.sqrt(-scores)
print("CV RMSLE:", rmsle_scores.mean(), "+/-", rmsle_scores.std())
```

CV RMSLE: 0.02791422699013636 +/- 0.004894436820739914

Lets find the feature importance of the model features


```
In [50]: model = RandomForestRegressor().fit(X_train, y_train)
importances = model.feature_importances_
sorted(zip(importances, X.columns), reverse=True)
```

```
Out [50]: [(np.float64(0.942884147192858), 'carat'),
(np.float64(0.033545345482517704), 'clarity'),
(np.float64(0.01488669080289693), 'color'),
(np.float64(0.004459897990771782), 'depth'),
(np.float64(0.0020440599977183437), 'table'),
(np.float64(0.0009222934705533204), 'Ideal'),
(np.float64(0.0005232910510988802), 'Premium'),
(np.float64(0.000477369391743848), 'Very Good'),
(np.float64(0.0002569046198413274), 'Good')]
```

Carat is the most important and dominant feature.

Now lets plot the Actuals vs predicted

```
In [51]: y_pred_rf = model.predict(X_test)
plt.scatter(y_test, model.predict(X_test), alpha=0.3)
plt.xlabel("Actual log(price)")
plt.ylabel("Predicted log(price)")
plt.title("Actual vs Predicted")
min_val = min(y_test.min(), y_pred_rf.min())
max_val = max(y_test.max(), y_pred_rf.max())
plt.plot([min_val, max_val], [min_val, max_val], 'r--', label='Perfect Predi
plt.show()
```

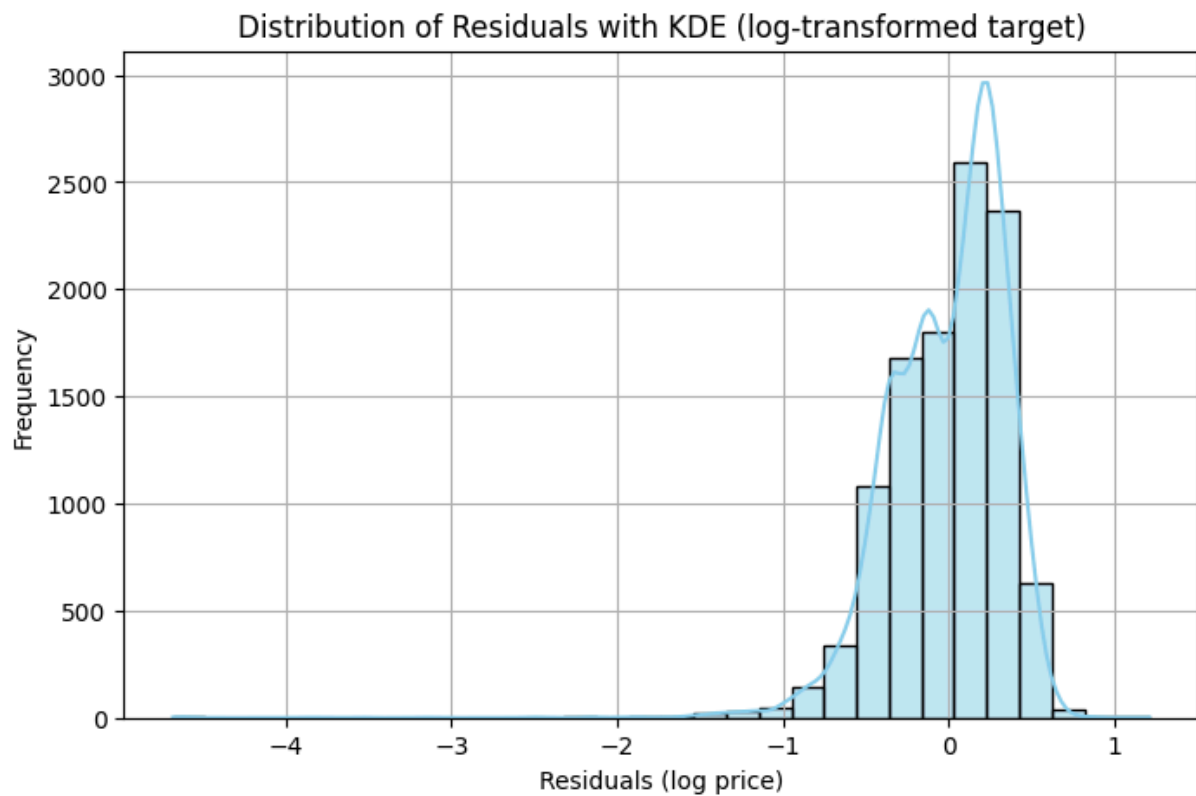


The Residual

- Looks roughly normal

```
In [52]: residuals = y_test - y_pred

# Plot
plt.figure(figsize=(8, 5))
sns.histplot(residuals, kde=True, bins=30, color='skyblue', edgecolor='black')
plt.title("Distribution of Residuals with KDE (log-transformed target)")
plt.xlabel("Residuals (log price)")
plt.ylabel("Frequency")
plt.grid(True)
plt.show()
```



Important! learn about VIF

Next steps

- Hyper parameter tuning using GridSearchCV or RandomizedSearchCV
- Change train test split and test